

# IVÁN CÓRDOVA

 [sebastian.cordova@pucp.edu.pe](mailto:sebastian.cordova@pucp.edu.pe)  [linkedin.com/in/ivan-cordova](https://linkedin.com/in/ivan-cordova)

## Education

### Pontifical Catholic University of Peru

BEng in Informatics Engineering, Minor in Artificial Intelligence

Lima, Peru

- Award: Best Bachelor Thesis for “Implementation of the Whale Optimization Algorithm in Financial Processes”.

## Experience

### Indra

AI Technical Lead

Lima, Peru

Feb 2025 - Current

- Architected and shipped an on-prem document QA pipeline for Peru’s public registry (SUNARP) processing 1M+ images/day using Python, YOLOv11, and OpenVINO behind FastAPI; achieved 80% mAP@0.5 and 12 ms/image on CPU via batching, replacing a manual QA workflow handled by a 15-person team.
- Built and launched a multi-tenant chatbot platform on AWS (FastAPI + Next.js) with agentic RAG and DSPy prompt optimization, delivering less than 500 ms p95 latency and 96% task accuracy; onboarded 5 internal teams and iterated using Logfire telemetry and customer feedback.
- Designed and implemented an AI recruiting platform for Indra, working end-to-end across backend (FastAPI, serverless AWS) and frontend (React); onboarded 500k+ CVs and used hybrid search with cross-encoder reranking to surface top candidates for Indra’s hiring pipeline. OCR scaled with Modal infrastructure using A100 GPUs and batch optimization.
- Led a 5-engineer squad, shipping 3 enterprise projects in 9 months and delivering 10+ AI PoCs (typically 2–3 days each) for leading banks and insurance companies in Peru and Latin America, establishing a fast iteration and experimentation culture.

### Astral Insights

Senior AI Engineer

Raleigh Nc, USA

Apr 2024 – Dec 2024

- Shipped an Azure-native Copilot for Power BI that turns natural-language questions into DAX/SQL using GPT-4o and meta-prompting; reduced cost per 1k queries by 65% vs Microsoft Fabric Copilot and contributed to closing a \$500k contract.
- Launched a LangGraph-based real-time supply-chain agent connected to package-tracking APIs and internal tools, improving incident detection and reducing downtime by 22% across 25 customers.
- Delivered 5+ LLM-powered PoCs in 2–3 day sprints for enterprise clients including Pfizer, GE Vernova, and Duke University, validating AI product ideas with real stakeholders.

### Yavendio (Startup)

Senior AI Engineer

Lima, Peru

Feb 2024 - Apr 2024

- Shipped a multi-agent vendor platform on Azure (Python, FastAPI, GPT-4o, LangGraph) for procurement, negotiation, and quoting, taking the product from concept to production in 4 weeks; this MVP became the foundation for raising \$900k in venture funding.
- Designed and implemented multimodal agentic RAG across catalogs, specs, and invoices (text + images) using hybrid retrieval (BM25 + pgvector), metadata filters, and cross-encoder reranking; embedding versioning and warm caches cut p95 latency by 70%.
- Delivered secure NL-to-SQL workflows that allowed non-technical users to generate analytics over production data while preserving governance, permissions, and auditability.

### Tuxpas

Software Engineer

Lima, Peru

Aug 2023 - Feb 2024

- Built and shipped an AWS-based data lake processing 1M+ medical records using Python ETL on S3, Glue, and Lambda orchestrated via EventBridge + Step Functions, with infra defined in AWS CDK, enabling new analytics products for clinicians and reducing cloud OPEX by 40%.

### BBVA

Software Analyst

Lima, Peru

Jan 2023 - Apr 2023

- Mapped and documented 15+ core financial processes at BBVA (business process diagrams and data flows), providing the foundation for later automation and efficiency improvements by engineering teams.

### Business Execution & Transformation Intern

Mar 2021 - Dic 2022

- Shipped 30+ features for internal finance tools using C, AutoHotkey, and JavaScript, translating manual system workflows into automated scripts used by 100 employees and finance staff, improving audit efficiency by 40%.

## Research & Teaching

---

### AI-PUCP (AI Research Group, Pontifical Catholic University of Peru)

Lima, Peru

Research Assistant

Jan 2025 – Jul 2025

- Explored Reinforcement Learning, Causal ML, and small language models through a series of research notebooks and experiments, focusing on practical applications and reproducible pipelines.
- Implemented and evaluated prototype models on curated datasets, documenting results and limitations to inform future lab projects and thesis directions.

### Pontifical Catholic University of Peru

Lima, Peru

Teaching Assistant — Advanced C++ Development

Feb 2024 – Apr 2024

- Supported a 50+ student course on Advanced C++ (data structures, algorithms, and modern C++ practices), resolving technical questions during labs and office hours.
- Reviewed assignments and projects, providing feedback on code quality, performance, and problem-solving approaches.

## Language Proficiency

---

Spanish: Native

English: Professional Proficiency (C1)

French: Basic (B1)

## Technical Skills

---

**Languages:** Python, Typescript, C, C++

**Developer Tools:** Linux, Docker, AWS, Azure, Cursor, Claude Code

**Technologies/Frameworks:** FastAPI, Next.js, React.js, Langchain, Langgraph, Dspy, Pytorch, vLLM

**Databases:** SQL, MongoDB, Redis, Qdrant, Pinecone, Neo4j